

DATA CLEANING

August 7, 2023

0.0.1 Finding missing data

```
[ ]:
```

```
[1]: import pandas as pd
import numpy as np
```

```
[ ]:
```

```
[2]: build_permits = pd.read_csv('Building_Permits.csv')
```

```
/usr/lib/python3/dist-packages/IPython/core/interactiveshell.py:3457:
DtypeWarning: Columns (22,32) have mixed types.Specify dtype option on import or
set low_memory=False.
  exec(code_obj, self.user_global_ns, self.user_ns)
```

```
[3]: build_permits.head()
```

```
[3]:
```

	Permit Number	Permit Type	Permit Type Definition	\
0	201505065519	4	sign - erect	
1	201604195146	4	sign - erect	
2	201605278609	3	additions alterations or repairs	
3	201611072166	8	otc alterations permit	
4	201611283529	6	demolitions	

	Permit	Creation Date	Block	Lot	Street Number	Street Number	Suffix	\
0		05/06/2015	0326	023	140		NaN	
1		04/19/2016	0306	007	440		NaN	
2		05/27/2016	0595	203	1647		NaN	
3		11/07/2016	0156	011	1230		NaN	
4		11/28/2016	0342	001	950		NaN	

	Street Name	Street Suffix	...	Existing Construction	Type	\
0	Ellis	St	...		3.0	
1	Geary	St	...		3.0	
2	Pacific	Av	...		1.0	
3	Pacific	Av	...		5.0	
4	Market	St	...		3.0	

	Existing Construction Type Description	Proposed Construction Type	\
0	constr type 3	NaN	
1	constr type 3	NaN	
2	constr type 1	1.0	
3	wood frame (5)	5.0	
4	constr type 3	NaN	

	Proposed Construction Type Description	Site Permit Supervisor	District	\
0	NaN	NaN	3.0	
1	NaN	NaN	3.0	
2	constr type 1	NaN	3.0	
3	wood frame (5)	NaN	3.0	
4	NaN	NaN	6.0	

	Neighborhoods - Analysis Boundaries	Zipcode	\
0	Tenderloin	94102.0	
1	Tenderloin	94102.0	
2	Russian Hill	94109.0	
3	Nob Hill	94109.0	
4	Tenderloin	94102.0	

	Location	Record ID
0	(37.785719256680785, -122.40852313194863)	1380611233945
1	(37.78733980600732, -122.41063199757738)	1420164406718
2	(37.7946573324287, -122.42232562979227)	1424856504716
3	(37.79595867909168, -122.41557405519474)	1443574295566
4	(37.78315261897309, -122.40950883997789)	144548169992

[5 rows x 43 columns]

```
[ ]:
```

```
[4]: build_permits.shape
```

```
[4]: (198900, 43)
```

```
[5]: build_permits.isnull().head()
```

	Permit Number	Permit Type	Permit Type Definition	Permit Creation Date	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

Block	Lot	Street Number	Street Number Suffix	Street Name	\
-------	-----	---------------	----------------------	-------------	---

0	False	False	False	True	False
1	False	False	False	True	False
2	False	False	False	True	False
3	False	False	False	True	False
4	False	False	False	True	False

	Street	Suffix	...	Existing Construction Type	\
0		False	...	False	
1		False	...	False	
2		False	...	False	
3		False	...	False	
4		False	...	False	

	Existing Construction Type	Description	Proposed Construction Type	\
0		False	True	
1		False	True	
2		False	False	
3		False	False	
4		False	True	

	Proposed Construction Type	Description	Site Permit	Supervisor	District	\
0		True	True		False	
1		True	True		False	
2		False	True		False	
3		False	True		False	
4		True	True		False	

	Neighborhoods - Analysis	Boundaries	Zipcode	Location	Record ID
0		False	False	False	False
1		False	False	False	False
2		False	False	False	False
3		False	False	False	False
4		False	False	False	False

[5 rows x 43 columns]

```
[6]: # finding total missing values
```

```
total_missing_count = build_permits.isnull().sum()
total_missing = total_missing_count.sum()
print(total_missing)
```

2245941

```
[ ]:
```

```
[7]: total_cells = np.product(build_permits.shape)
total_cells
```

```
[7]: 8552700
```

```
[8]: # Percent of data that is missing

percent_missing = (total_missing/total_cells) * 100
print('Percent of data that is missing : ',percent_missing)
```

```
Percent of data that is missing : 26.26002315058403
```

```
[9]: build_permits.shape
```

```
[9]: (198900, 43)
```

0.0.2 Drop missing Values

```
[10]: build_permits.dropna()
```

```
[10]: Empty DataFrame
Columns: [Permit Number, Permit Type, Permit Type Definition, Permit Creation
Date, Block, Lot, Street Number, Street Number Suffix, Street Name, Street
Suffix, Unit, Unit Suffix, Description, Current Status, Current Status Date,
Filed Date, Issued Date, Completed Date, First Construction Document Date,
Structural Notification, Number of Existing Stories, Number of Proposed Stories,
Voluntary Soft-Story Retrofit, Fire Only Permit, Permit Expiration Date,
Estimated Cost, Revised Cost, Existing Use, Existing Units, Proposed Use,
Proposed Units, Plansets, TIDF Compliance, Existing Construction Type, Existing
Construction Type Description, Proposed Construction Type, Proposed Construction
Type Description, Site Permit, Supervisor District, Neighborhoods - Analysis
Boundaries, Zipcode, Location, Record ID]
Index: []
```

```
[0 rows x 43 columns]
```

```
[11]: build_permits.shape
```

```
[11]: (198900, 43)
```

```
[12]: # Remove the column with atleast one missing value

drop_column_wise = build_permits.dropna(axis = 1)
drop_column_wise.head()
```

```
[12]:  Permit Number  Permit Type  Permit Type Definition \
0  201505065519          4          sign - erect
```

1	201604195146	4	sign - erect
2	201605278609	3	additions alterations or repairs
3	201611072166	8	otc alterations permit
4	201611283529	6	demolitions

	Permit	Creation Date	Block	Lot	Street Number	Street Name	Current Status	\
0		05/06/2015	0326	023	140	Ellis	expired	
1		04/19/2016	0306	007	440	Geary	issued	
2		05/27/2016	0595	203	1647	Pacific	withdrawn	
3		11/07/2016	0156	011	1230	Pacific	complete	
4		11/28/2016	0342	001	950	Market	issued	

	Current Status	Date	Filed Date	Record ID
0		12/21/2017	05/06/2015	1380611233945
1		08/03/2017	04/19/2016	1420164406718
2		09/26/2017	05/27/2016	1424856504716
3		07/24/2017	11/07/2016	1443574295566
4		12/01/2017	11/28/2016	144548169992

```
[13]: drop_column_wise.shape
```

```
[13]: (198900, 12)
```

```
[14]: # finding the how much data we lost?

#
original_data = build_permits.shape[1]
#column wise we lose
col_lose = drop_column_wise.shape[1]

#how much data we lose
data_lose = original_data - col_lose
data_lose
```

```
[14]: 31
```

```
[ ]:
```

0.1 Filling in Missing Values Automatically

```
[ ]:
```

```
[15]: build_permits.fillna(0).head(2)
```

```
[15]: Permit Number  Permit Type  Permit Type Definition  Permit Creation Date  \
0    201505065519          4          sign - erect          05/06/2015
1    201604195146          4          sign - erect          04/19/2016
```

	Block	Lot	Street Number	Street Number	Suffix	Street Name	Street Suffix	\
0	0326	023	140		0	Ellis	St	
1	0306	007	440		0	Geary	St	

	...	Existing Construction Type	Existing Construction Type	Description	\
0	...	3.0		constr type 3	
1	...	3.0		constr type 3	

	Proposed Construction Type	Proposed Construction Type	Description	\
0	0.0		0	
1	0.0		0	

	Site Permit Supervisor District	Neighborhoods - Analysis	Boundaries	\
0	0	3.0	Tenderloin	
1	0	3.0	Tenderloin	

	Zipcode	Location	Record ID
0	94102.0	(37.785719256680785, -122.40852313194863)	1380611233945
1	94102.0	(37.78733980600732, -122.41063199757738)	1420164406718

[2 rows x 43 columns]

```
[16]: build_permits.head(2)
```

```
[16]: Permit Number Permit Type Permit Type Definition Permit Creation Date \
0 201505065519 4 sign - erect 05/06/2015
1 201604195146 4 sign - erect 04/19/2016
```

	Block	Lot	Street Number	Street Number	Suffix	Street Name	Street Suffix	\
0	0326	023	140		NaN	Ellis	St	
1	0306	007	440		NaN	Geary	St	

	...	Existing Construction Type	Existing Construction Type	Description	\
0	...	3.0		constr type 3	
1	...	3.0		constr type 3	

	Proposed Construction Type	Proposed Construction Type	Description	\
0	NaN		NaN	
1	NaN		NaN	

	Site Permit Supervisor District	Neighborhoods - Analysis	Boundaries	\
0	NaN	3.0	Tenderloin	
1	NaN	3.0	Tenderloin	

	Zipcode	Location	Record ID
0	94102.0	(37.785719256680785, -122.40852313194863)	1380611233945

```
1 94102.0 (37.78733980600732, -122.41063199757738) 1420164406718
```

```
[2 rows x 43 columns]
```

```
[23]: # replace all NAN's the values that comes directly after/before it in the same
      ↪column,
      #then replace all the remaining NAN's with 0

      fill_null = build_permits.fillna(method = 'bfill',axis = 1, inplace = False).
      ↪fillna(0).head()
```

```
[24]: fill_null
```

```
[24]: Permit Number Permit Type Permit Type Definition \
0 201505065519 4 sign - erect
1 201604195146 4 sign - erect
2 201605278609 3 additions alterations or repairs
3 201611072166 8 otc alterations permit
4 201611283529 6 demolitions
```

```
Permit Creation Date Block Lot Street Number Street Number Suffix \
0 05/06/2015 0326 023 140 Ellis
1 04/19/2016 0306 007 440 Geary
2 05/27/2016 0595 203 1647 Pacific
3 11/07/2016 0156 011 1230 Pacific
4 11/28/2016 0342 001 950 Market
```

```
Street Name Street Suffix ... Existing Construction Type \
0 Ellis St ... 3.0
1 Geary St ... 3.0
2 Pacific Av ... 1.0
3 Pacific Av ... 5.0
4 Market St ... 3.0
```

```
Existing Construction Type Description Proposed Construction Type \
0 constr type 3 3.0
1 constr type 3 3.0
2 constr type 1 1.0
3 wood frame (5) 5.0
4 constr type 3 6.0
```

```
Proposed Construction Type Description Site Permit Supervisor District \
0 3.0 3.0 3.0
1 3.0 3.0 3.0
2 constr type 1 3.0 3.0
3 wood frame (5) 3.0 3.0
4 6.0 6.0 6.0
```

	Neighborhoods - Analysis Boundaries	Zipcode	\
0	Tenderloin	94102.0	
1	Tenderloin	94102.0	
2	Russian Hill	94109.0	
3	Nob Hill	94109.0	
4	Tenderloin	94102.0	

	Location	Record ID
0	(37.785719256680785, -122.40852313194863)	1380611233945
1	(37.78733980600732, -122.41063199757738)	1420164406718
2	(37.7946573324287, -122.42232562979227)	1424856504716
3	(37.79595867909168, -122.41557405519474)	1443574295566
4	(37.78315261897309, -122.40950883997789)	144548169992

[5 rows x 43 columns]

```
[19]: check_isnull = fill_null.isnull().sum()
      check_isnull[10:30]
```

```
[19]: Unit                                0
      Unit Suffix                        0
      Description                        0
      Current Status                     0
      Current Status Date                 0
      Filed Date                         0
      Issued Date                        0
      Completed Date                     0
      First Construction Document Date    0
      Structural Notification              0
      Number of Existing Stories           0
      Number of Proposed Stories           0
      Voluntary Soft-Story Retrofit        0
      Fire Only Permit                    0
      Permit Expiration Date               0
      Estimated Cost                       0
      Revised Cost                        0
      Existing Use                         0
      Existing Units                       0
      Proposed Use                         0
      dtype: int64
```

```
[20]: # Save the dataset after modifications

      #fill_null.to_csv('Building_Permits.csv')
```

```
[74]: #df = pd.read_csv('Building_Permits.csv')
```



```
[75]: #df = df.drop('Unnamed: 0',axis = 1)
```

```
[76]: #df.isnull().sum()[10:30]
```

```
[77]: #df.describe()
```

```
[28]: #df.describe().transpose
```

```
[78]: #df.head()
```

0.2 Scaling And Normalization

```
[30]: #!pip install stats
```

```
[31]: #!pip install scipy
```

```
[32]: import pandas as pd
import numpy as np
#Plotting modules
import seaborn as sns
import matplotlib.pyplot as plt
#for box-cox transformation
from scipy import stats
# for minmax scaling
from mlxtend.preprocessing import minmax_scaling
# set seed for reproducibility
np.random.seed(0)
```

```
/usr/lib/python3/dist-packages/scipy/__init__.py:146: UserWarning: A NumPy
version >=1.17.3 and <1.25.0 is required for this version of SciPy (detected
version 1.25.2
```

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

0.3 Scaling

This means that you're transforming your data so that it fits within a specific scale, like 0-100 or 0-1. 1) You want to scale data when you're using methods based on measures of how far apart data points are, like support vector machines (SVM) or k-nearest neighbors (KNN). With these algorithms, a change of "1" in any numeric feature is given the same importance.

```
[33]: # Generate 1000 data points randomly drawn from an exponential distribution
```

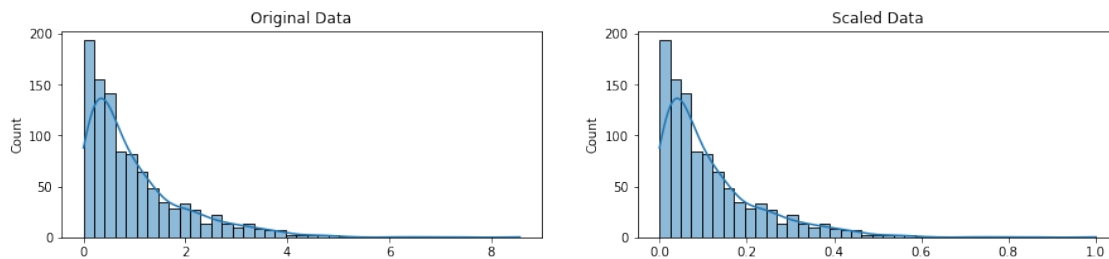
```
original_data = np.random.exponential(size = 1000)
#original_data

# Min-max scale the data b/w 0 & 1
scaled_data = minmax_scaling(original_data,columns=[0])
```

```
#plot both together to compare

fig, ax = plt.subplots(1,2,figsize = (15,3))
#Original plot
sns.histplot(original_data, ax = ax[0], kde = True, legend= False)
ax[0].set_title("Original Data")
#scaled plot
sns.histplot(scaled_data,ax = ax[1], kde = True, legend= False)
ax[1].set_title("Scaled Data")

plt.show()
```

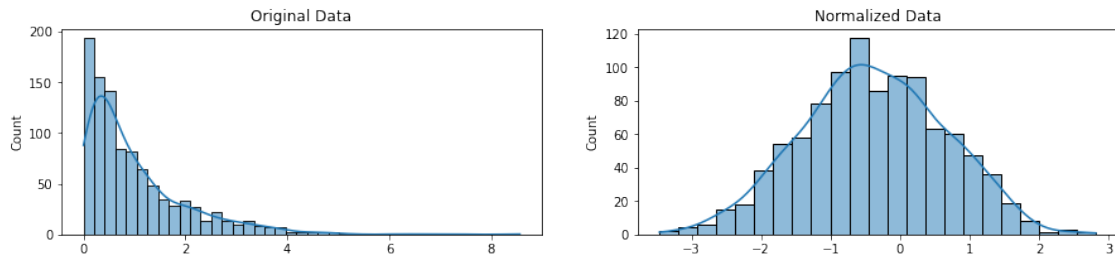


0.4 Normalization

- 1) Scaling just changes the range of your data. Normalization is a more radical transformation. The point of normalization is to change your observations so that they can be described as a normal distribution.
- 2) Normal distribution: Also known as the “bell curve”, this is a specific statistical distribution where a roughly equal observations fall above and below the mean, the mean and the median are the same, and there are more observations closer to the mean. The normal distribution is also known as the Gaussian distribution.
- 3) In general, you’ll normalize your data if you’re going to be using a machine learning or statistics technique that assumes your data is normally distributed. Some examples of these include **linear discriminant analysis (LDA)** and **Gaussian naïve Bayes**. (Pro tip: any method with “Gaussian” in the name probably assumes normality.)
- 4) The method we’re using to normalize here is called the **Box-Cox Transformation**.

```
[34]: normalized_data = stats.boxcox(original_data)
# plot both together to compare
fig, ax = plt.subplots(1,2,figsize = (15,3))
sns.histplot(original_data,ax = ax[0], kde = True,legend=False)
ax[0].set_title("Original Data")
sns.histplot(normalized_data[0],ax = ax[1],kde = True,legend=False)
ax[1].set_title("Normalized Data")
```

```
plt.show()
```



```
[35]: kickstarters_2017 = pd.read_csv("ks-projects-201801.csv")
```

```
[36]: kickstarters_2017.head()
```

```
[36]:
```

	ID	name \
0	1000002330	The Songs of Adelaide & Abullah
1	1000003930	Greeting From Earth: ZGAC Arts Capsule For ET
2	1000004038	Where is Hank?
3	1000007540	ToshiCapital Rekordz Needs Help to Complete Album
4	1000011046	Community Film Project: The Art of Neighborhoo...

	category	main_category	currency	deadline	goal \
0	Poetry	Publishing	GBP	2015-10-09	1000.0
1	Narrative Film	Film & Video	USD	2017-11-01	30000.0
2	Narrative Film	Film & Video	USD	2013-02-26	45000.0
3	Music	Music	USD	2012-04-16	5000.0
4	Film & Video	Film & Video	USD	2015-08-29	19500.0

	launched	pledged	state	backers	country	usd pledged \
0	2015-08-11 12:12:28	0.0	failed	0	GB	0.0
1	2017-09-02 04:43:57	2421.0	failed	15	US	100.0
2	2013-01-12 00:20:50	220.0	failed	3	US	220.0
3	2012-03-17 03:24:11	1.0	failed	1	US	1.0
4	2015-07-04 08:35:03	1283.0	canceled	14	US	1283.0

	usd_pledged_real	usd_goal_real
0	0.0	1533.95
1	2421.0	30000.00
2	220.0	45000.00
3	1.0	5000.00
4	1283.0	19500.00

```
[37]: # select the usd_goal_real column
original_data = pd.DataFrame(kickstarters_2017.usd_goal_real)
```

```

# scale the goals from 0 to 1
scaled_data = minmax_scaling(original_data, columns=['usd_goal_real'])

print('Original data\nPreview:\n', original_data.head())
print('Minimum value:', float(original_data.min()),
      '\nMaximum value:', float(original_data.max()))
print('_'*30)

print('\nScaled data\nPreview:\n', scaled_data.head())
print('Minimum value:', float(scaled_data.min()),
      '\nMaximum value:', float(scaled_data.max()))

```

Original data

Preview:

	usd_goal_real
0	1533.95
1	30000.00
2	45000.00
3	5000.00
4	19500.00

Minimum value: 0.01

Maximum value: 166361390.71

Scaled data

Preview:

	usd_goal_real
0	0.000009
1	0.000180
2	0.000270
3	0.000030
4	0.000117

Minimum value: 0.0

Maximum value: 1.0

```

[38]: original_goal_data = pd.DataFrame(kickstarters_2017.goal)
      original_goal_data.head(6)

```

```

[38]:      goal
0    1000.0
1   30000.0
2   45000.0
3    5000.0
4   19500.0
5   50000.0

```

```
[39]: scaled_goal_data = minmax_scaling(original_goal_data, columns=['goal'])
scaled_goal_data.head(6)
```

```
[39]:      goal
0  0.000010
1  0.000300
2  0.000450
3  0.000050
4  0.000195
5  0.000500
```

```
[40]: # #scaled plot
# sns.histplot(scaled_goal_data, kde = True, legend= False)
# plt.title("Scaled Data")

# plt.show()
```

```
[41]: # get the index of all positive pledges (Box-Cox only takes positive values)
index_of_positive_pledges = kickstarters_2017.usd_pledged_real > 0

# get only positive pledges (using their indexes)
positive_pledges = kickstarters_2017.usd_pledged_real.
    ↪loc[index_of_positive_pledges]

# normalize the pledges (w/ Box-Cox)
normalized_pledges = pd.Series(stats.boxcox(positive_pledges)[0],
                                name='usd_pledged_real', index=positive_pledges.
    ↪index)

print('Original data\nPreview:\n', positive_pledges.head())
print('Minimum value:', float(positive_pledges.min()),
      '\nMaximum value:', float(positive_pledges.max()))
print('_'*30)

print('\nNormalized data\nPreview:\n', normalized_pledges.head())
print('Minimum value:', float(normalized_pledges.min()),
      '\nMaximum value:', float(normalized_pledges.max()))
```

Original data

Preview:

```
1      2421.0
2      220.0
3         1.0
4     1283.0
5    52375.0
```

Name: usd_pledged_real, dtype: float64

Minimum value: 0.45

Maximum value: 20338986.27

Normalized data

Preview:

```
1    10.165142
2     6.468598
3     0.000000
4     9.129277
5    15.836853
```

Name: `usd_pledged_real`, dtype: float64

Minimum value: -0.7779954122762203

Maximum value: 30.69054020451361

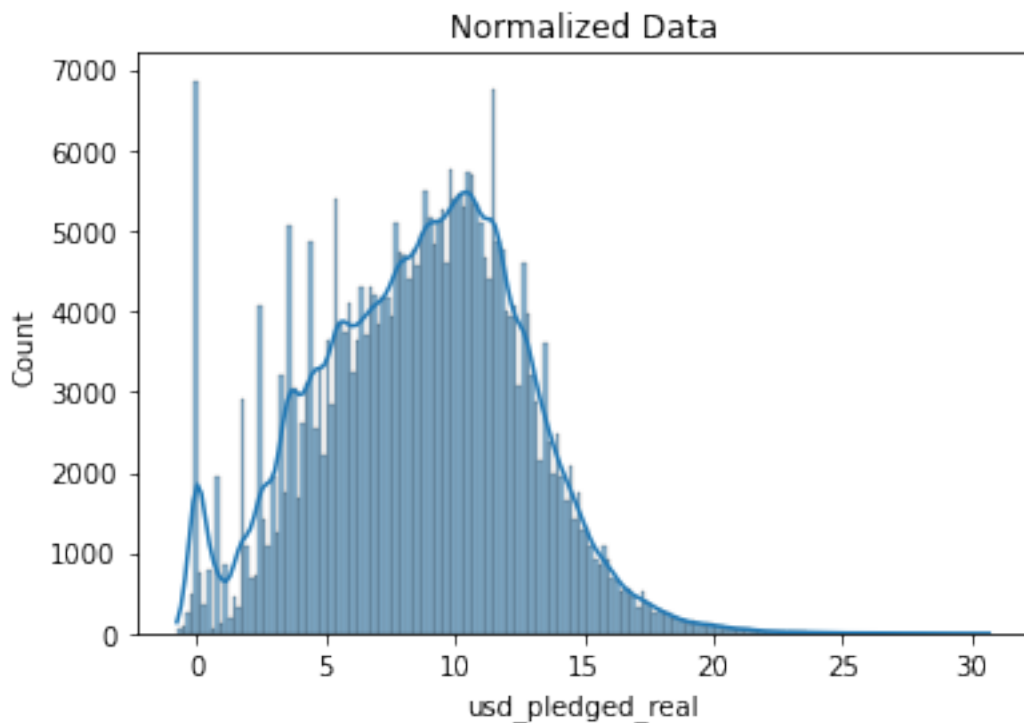
```
[73]: # # get the index of all positive pledges (Box-Cox only takes positive values)
      # index_of_positive_pledges = kickstarters_2017.usd_pledged_real > 0

      # # get only positive pledges (using their indexes)
      # positive_pledges = kickstarters_2017.usd_pledged_real.
      #     ↪ loc[index_of_positive_pledges]

      # # normalize the pledges (w/ Box-Cox)
      # normalized_pledges = pd.Series(stats.boxcox(positive_pledges)[0],
      #     ↪                                name='usd_pledged_real',
      #     ↪ index=positive_pledges.index)

      sns.histplot(normalized_pledges, kde = True, legend= False)
      plt.title("Normalized Data")

      plt.show()
```



```
[ ]:
```

0.5 Parsing Dates

```
[43]: earthquakes = pd.read_csv('Earthquake_database.csv')
```

```
[44]: earthquakes.head()
```

```
[44]:
```

	Date	Time	Latitude	Longitude	Type	Depth	Depth Error	\
0	01/02/1965	13:44:18	19.246	145.616	Earthquake	131.6	NaN	
1	01/04/1965	11:29:49	1.863	127.352	Earthquake	80.0	NaN	
2	01/05/1965	18:05:58	-20.579	-173.972	Earthquake	20.0	NaN	
3	01/08/1965	18:49:43	-59.076	-23.557	Earthquake	15.0	NaN	
4	01/09/1965	13:32:50	11.938	126.427	Earthquake	15.0	NaN	

	Depth	Seismic Stations	Magnitude	Magnitude	Type	...	\
0		NaN	6.0		MW	...	
1		NaN	5.8		MW	...	
2		NaN	6.2		MW	...	
3		NaN	5.8		MW	...	
4		NaN	5.8		MW	...	

	Magnitude	Seismic Stations	Azimuthal Gap	Horizontal Distance	\
--	-----------	------------------	---------------	---------------------	---

0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

	Horizontal Error	Root Mean Square	ID	Source Location	Source \
0	NaN	NaN	ISCGEM860706	ISCGEM	ISCGEM
1	NaN	NaN	ISCGEM860737	ISCGEM	ISCGEM
2	NaN	NaN	ISCGEM860762	ISCGEM	ISCGEM
3	NaN	NaN	ISCGEM860856	ISCGEM	ISCGEM
4	NaN	NaN	ISCGEM860890	ISCGEM	ISCGEM

	Magnitude Source	Status
0	ISCGEM	Automatic
1	ISCGEM	Automatic
2	ISCGEM	Automatic
3	ISCGEM	Automatic
4	ISCGEM	Automatic

[5 rows x 21 columns]

```
[45]: earthquakes.isnull().sum()
```

```
[45]: Date                0
      Time                0
      Latitude            0
      Longitude           0
      Type                0
      Depth               0
      Depth Error         18951
      Depth Seismic Stations 16315
      Magnitude           0
      Magnitude Type       3
      Magnitude Error      23085
      Magnitude Seismic Stations 20848
      Azimuthal Gap        16113
      Horizontal Distance   21808
      Horizontal Error      22256
      Root Mean Square     6060
      ID                  0
      Source              0
      Location Source      0
      Magnitude Source     0
      Status              0
      dtype: int64
```



```
[46]: earthquakes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23412 entries, 0 to 23411
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                23412 non-null  object
1   Time                                23412 non-null  object
2   Latitude                            23412 non-null  float64
3   Longitude                           23412 non-null  float64
4   Type                                23412 non-null  object
5   Depth                               23412 non-null  float64
6   Depth Error                         4461 non-null   float64
7   Depth Seismic Stations              7097 non-null   float64
8   Magnitude                           23412 non-null  float64
9   Magnitude Type                      23409 non-null  object
10  Magnitude Error                     327 non-null    float64
11  Magnitude Seismic Stations          2564 non-null   float64
12  Azimuthal Gap                       7299 non-null   float64
13  Horizontal Distance                 1604 non-null   float64
14  Horizontal Error                    1156 non-null   float64
15  Root Mean Square                   17352 non-null  float64
16  ID                                  23412 non-null  object
17  Source                             23412 non-null  object
18  Location Source                     23412 non-null  object
19  Magnitude Source                    23412 non-null  object
20  Status                             23412 non-null  object
dtypes: float64(12), object(9)
memory usage: 3.8+ MB
```

```
[47]: earthquakes['Date'].head()
```

```
[47]: 0    01/02/1965
1    01/04/1965
2    01/05/1965
3    01/08/1965
4    01/09/1965
Name: Date, dtype: object
```

```
[ ]:
```

- 1) You may have to check the numpy documentation to match the letter code to the dtype of the object. “O” is the code for “object”, so we can see that these two methods give us the same information.

```
[48]: earthquakes['Date'].dtype
```

```
[48]: dtype('O')
```

0.5.1 Convert our date columns to datetime

- 1) Now that we know that our date column isn't being recognized as a date, it's time to convert it so that it is recognized as a date. This is called "parsing dates" because we're taking in a string and identifying its component parts.
- 2) We can determine what the format of our dates are with a guide called "strftime directive", which you can find more information on at this link. The basic idea is that you need to point out which parts of the date are where and what punctuation is between them. There are lots of possible parts of a date, but the most common are %d for day, %m for month, %y for a two-digit year and %Y for a four digit year.
- 3) Some examples:

- 1/17/07 has the format "%m/%d/%y"
- 17-1-2007 has the format "%d-%m-%Y"

- 4) Looking back up at the head of the "date" column in the landslides dataset, we can see that it's in the format "month/day/two-digit year", so we can use the same syntax as the first example to parse in our dates:

```
[ ]:
```

```
[49]: date_length = earthquakes.Date.str.len()
      date_length.value_counts()
```

```
[49]: 10    23409
      24      3
      Name: Date, dtype: int64
```

```
[50]: indices = np.where([date_length == 24])[1]
      print('Indices with currepted data : ',indices)
      earthquakes.loc[indices]
```

```
Indices with currepted data : [ 3378  7512 20650]
```

```
[50]:
```

	Date	Time	Latitude	\
3378	1975-02-23T02:58:41.000Z	1975-02-23T02:58:41.000Z	8.017	
7512	1985-04-28T02:53:41.530Z	1985-04-28T02:53:41.530Z	-32.998	
20650	2011-03-13T02:23:34.520Z	2011-03-13T02:23:34.520Z	36.344	

	Longitude	Type	Depth	Depth Error	Depth Seismic Stations	\
3378	124.075	Earthquake	623.0	NaN	NaN	
7512	-71.766	Earthquake	33.0	NaN	NaN	
20650	142.344	Earthquake	10.1	13.9	289.0	

	Magnitude	Magnitude Type	...	Magnitude Seismic Stations	\
3378	5.6	MB	...	NaN	

7512	5.6	MW	...	NaN
20650	5.8	MWC	...	NaN

	Azimuthal Gap	Horizontal Distance	Horizontal Error	Root Mean Square	\
3378	NaN	NaN	NaN	NaN	
7512	NaN	NaN	NaN	1.30	
20650	32.3	NaN	NaN	1.06	

	ID	Source	Location	Source	Magnitude	Source	Status
3378	USP0000A09	US		US		US	Reviewed
7512	USP0002E81	US		US		HRV	Reviewed
20650	USP000HWQP	US		US		GCMT	Reviewed

[3 rows x 21 columns]

```
[51]: earthquakes.loc[3378, 'Date'] = '02/23/1975'
earthquakes.loc[7512, 'Date'] = '04/28/1985'
earthquakes.loc[20650, 'Date'] = '03/13/2011'

earthquakes['date_parsed'] = pd.to_datetime(earthquakes['Date'], format='%m/%d/%Y')
```

- 1) What if I run into an error with multiple date formats? While we're specifying the date format here, sometimes you'll run into an error when there are multiple date formats in a single column. If that happens, you can have pandas try to infer what the right date format should be. You can do that like so: `landslides['date_parsed']=pd.to_datetime(landslides['Date'],infer_datetime_format=True)`
- 2) Why don't you always use `'infer_datetime_format = True'`? There are two big reasons not to always have pandas guess the time format. The first is that pandas won't always be able to figure out the correct date format, especially if someone has gotten creative with data entry. The second is that it's much slower than specifying the exact format of the dates.

[]:

```
[52]: earthquakes.date_parsed.head()
```

```
[52]: 0    1965-01-02
1    1965-01-04
2    1965-01-05
3    1965-01-08
4    1965-01-09
Name: date_parsed, dtype: datetime64[ns]
```

```
[53]: earthquakes.head()
```

```
[53]:
```

	Date	Time	Latitude	Longitude	Type	Depth	Depth Error	\
0	01/02/1965	13:44:18	19.246	145.616	Earthquake	131.6		NaN
1	01/04/1965	11:29:49	1.863	127.352	Earthquake	80.0		NaN
2	01/05/1965	18:05:58	-20.579	-173.972	Earthquake	20.0		NaN
3	01/08/1965	18:49:43	-59.076	-23.557	Earthquake	15.0		NaN
4	01/09/1965	13:32:50	11.938	126.427	Earthquake	15.0		NaN

	Depth	Seismic Stations	Magnitude	Magnitude Type	...	Azimuthal Gap	\
0		NaN	6.0	MW	...	NaN	
1		NaN	5.8	MW	...	NaN	
2		NaN	6.2	MW	...	NaN	
3		NaN	5.8	MW	...	NaN	
4		NaN	5.8	MW	...	NaN	

	Horizontal Distance	Horizontal Error	Root Mean Square	ID	\
0	NaN	NaN	NaN	ISCGEM860706	
1	NaN	NaN	NaN	ISCGEM860737	
2	NaN	NaN	NaN	ISCGEM860762	
3	NaN	NaN	NaN	ISCGEM860856	
4	NaN	NaN	NaN	ISCGEM860890	

	Source Location	Source Magnitude	Source	Status	date_parsed
0	ISCGEM	ISCGEM	ISCGEM	Automatic	1965-01-02
1	ISCGEM	ISCGEM	ISCGEM	Automatic	1965-01-04
2	ISCGEM	ISCGEM	ISCGEM	Automatic	1965-01-05
3	ISCGEM	ISCGEM	ISCGEM	Automatic	1965-01-08
4	ISCGEM	ISCGEM	ISCGEM	Automatic	1965-01-09

[5 rows x 22 columns]

0.5.2 Select the day of the month

- 1) Now that we have a column of parsed dates, we can extract information like the day of the month that a landslide occurred.

```
[ ]:
```

```
[54]: day_of_month = earthquakes['date_parsed'].dt.day
      day_of_month.head()
```

```
[54]: 0    2
      1    4
      2    5
      3    8
      4    9
      Name: date_parsed, dtype: int64
```

0.5.3 Plot the day of the month to check the date parsing

- 1) One of the biggest dangers in parsing dates is mixing up the months and days. The `to_datetime()` function does have very helpful error messages, but it doesn't hurt to double-check that the days of the month we've extracted make sense.
- 2) To do this, let's plot a histogram of the days of the month. We expect it to have values between 1 and 31 and, since there's no reason to suppose the landslides are more common on some days of the month than others, a relatively even distribution. (With a dip on 31 because not all months have 31 days.) Let's see if that's the case:

[]:

```
[55]: # remove na's
      day_of_month = day_of_month.dropna()
      #plot
      sns.distplot(day_of_month,kde = False,bins = 31)
      plt.show()
```

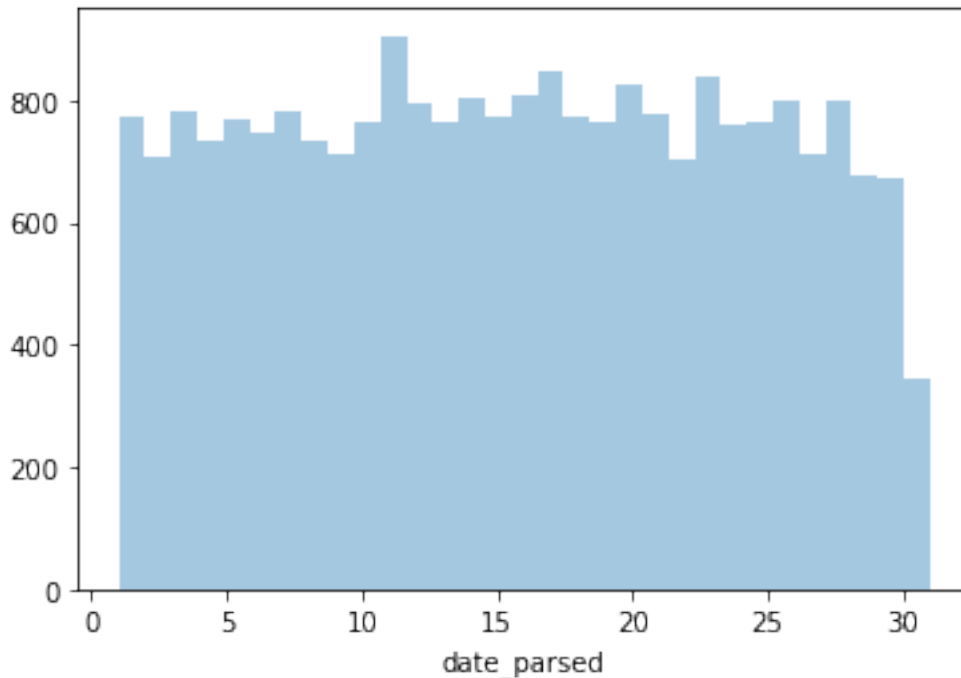
/tmp/ipykernel_5179/3349062656.py:4: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(day_of_month,kde = False,bins = 31)
```



0.6 Inconsistent data Entry

```
[56]: #!pip install fuzzywuzzy
```

```
[57]: #!pip install charset_normalizer
```

```
[58]: import pandas as pd
import numpy as np
#Helpful modules
import fuzzywuzzy
from fuzzywuzzy import process
import charset_normalizer
```

```
/home/mahesh/.local/lib/python3.10/site-packages/fuzzywuzzy/fuzz.py:11:
UserWarning: Using slow pure-python SequenceMatcher. Install python-Levenshtein
to remove this warning
  warnings.warn('Using slow pure-python SequenceMatcher. Install python-
Levenshtein to remove this warning')
```

```
[59]: professors = pd.read_csv('pakistan_intellectual_capital.csv')
```

```
[60]: np.random.seed(0)
```

```
[61]: professors = professors.drop('Unnamed: 0',axis = 1)
```

0.6.1 Do some preliminary text pre-processing

```
[ ]:
```

```
[62]: professors.head()
```

```
[62]:  S#          Teacher Name          University Currently Teaching \
0    3      Dr. Abdul Basit          University of Balochistan
1    5      Dr. Waheed Noor          University of Balochistan
2    6      Dr. Junaid Baber          University of Balochistan
3    7  Dr. Maheen Bakhtyar          University of Balochistan
4   25          Samina Azim  Sardar Bahadur Khan Women's University

          Department Province University Located          Designation \
0  Computer Science & IT          Balochistan  Assistant Professor
1  Computer Science & IT          Balochistan  Assistant Professor
2  Computer Science & IT          Balochistan  Assistant Professor
3  Computer Science & IT          Balochistan  Assistant Professor
4    Computer Science          Balochistan          Lecturer

Terminal Degree          Graduated from \
0          PhD          Asian Institute of Technology
1          PhD          Asian Institute of Technology
2          PhD          Asian Institute of Technology
3          PhD          Asian Institute of Technology
4          BS  Balochistan University of Information Technolo...

Country    Year          Area of Specialization/Research Interests \
0  Thailand    NaN          Software Engineering & DBMS
1  Thailand    NaN          DBMS
2  Thailand    NaN          Information processing, Multimedia mining
3  Thailand    NaN  NLP, Information Retrieval, Question Answering...
4  Pakistan    2005.0          VLSI Electronics DLD Database

Other Information
0          NaN
1          NaN
2          NaN
3          NaN
4          NaN
```

```
[63]: countries = professors['Country'].unique()
```

```
[64]: countries.sort()
countries
```

```
[64]: array([' Germany', ' New Zealand', ' Sweden', ' USA', 'Australia',
        'Austria', 'Canada', 'China', 'Finland', 'France', 'Greece',
        'HongKong', 'Ireland', 'Italy', 'Japan', 'Macau', 'Malaysia',
        'Mauritius', 'Netherland', 'New Zealand', 'Norway', 'Pakistan',
        'Portugal', 'Russian Federation', 'Saudi Arabia', 'Scotland',
        'Singapore', 'South Korea', 'SouthKorea', 'Spain', 'Sweden',
        'Thailand', 'Turkey', 'UK', 'USA', 'USofA', 'Urbana', 'germany'],
        dtype=object)
```

[]:

- 1) Just looking at this, I can see some problems due to inconsistent data entry: ' Germany', and 'germany', for example, or ' New Zealand' and 'New Zealand'.
- 2) The first thing I'm going to do is make everything lower case (I can change it back at the end if I like) and remove any white spaces at the beginning and end of cells. Inconsistencies in capitalizations and trailing white spaces are very common in text data and you can fix a good 80% of your text data entry inconsistencies by doing this.

[]:

```
[65]: professors['Country'] = professors['Country'].str.lower()
professors['Country'] = professors['Country'].str.strip()
```

```
[66]: countries = professors['Country'].unique()
countries.sort()
countries
```

```
[66]: array(['australia', 'austria', 'canada', 'china', 'finland', 'france',
        'germany', 'greece', 'hongkong', 'ireland', 'italy', 'japan',
        'macau', 'malaysia', 'mauritius', 'netherland', 'new zealand',
        'norway', 'pakistan', 'portugal', 'russian federation',
        'saudi arabia', 'scotland', 'singapore', 'south korea',
        'southkorea', 'spain', 'sweden', 'thailand', 'turkey', 'uk',
        'urbana', 'usa', 'usofa'], dtype=object)
```

[]:

0.6.2 Use fuzzy matching to correct inconsistent data entry

Alright, let's take another look at the 'Country' column and see if there's any more data cleaning we need to do.

[]:

It does look like there is another inconsistency: 'southkorea' and 'south korea' should be the same.

We're going to use the fuzzywuzzy package to help identify which strings are closest to each other. This dataset is small enough that we could probably correct errors by hand, but that approach

doesn't scale well. (Would you want to correct a thousand errors by hand? What about ten thousand? Automating things as early as possible is generally a good idea. Plus, it's fun!)

Fuzzy matching:

- 3) The process of automatically finding text strings that are very similar to the target string. In general, a string is considered "closer" to another one the fewer characters you'd need to change if you were transforming one string into another. So "apple" and "snapple" are two changes away from each other (add "s" and "n") while "in" and "on" and one change away (replace "i" with "o"). You won't always be able to rely on fuzzy matching 100%, but it will usually end up saving you at least a little time.

Fuzzywuzzy returns a ratio given two strings. The closer the ratio is to 100, the smaller the edit distance between the two strings. Here, we're going to get the ten strings from our list of cities that have the closest distance to "south korea".

```
[67]: matches = fuzzywuzzy.process.extract('south korea',countries,limit = 10,
                                           scorer=fuzzywuzzy.fuzz.token_sort_ratio)
```

```
[68]: matches
      #matches
```

```
[68]: [('south korea', 100),
      ('southkorea', 48),
      ('saudi arabia', 43),
      ('norway', 35),
      ('ireland', 33),
      ('portugal', 32),
      ('singapore', 30),
      ('netherland', 29),
      ('macau', 25),
      ('usofa', 25)]
```

```
[ ]:
```

We can see that two of the items in the cities are very close to "south korea": "south korea" and "southkorea". Let's replace all rows in our "Country" column that have a ratio of > 47 with "south korea".

To do this, I'm going to write a function. (It's a good idea to write a general purpose function you can reuse if you think you might have to do a specific task more than once or twice. This keeps you from having to copy and paste code too often, which saves time and can help prevent mistakes.)

```
[ ]:
```

```
[69]: def replace_matches_in_column(df,column,string_to_match,min_ratio = 47):
      #get a list of unique strings
      string = df[column].unique()
      #get the top 10 closest matches to our input string
```

```

matches = fuzzywuzzy.process.extract(string_to_match,string,limit = 10,
                                     scorer=fuzzywuzzy.fuzz.
↳token_sort_ratio)
    # only get matches with ratio > 90
    close_matches = [matches[0] for matches in matches if matches[1] >=
↳min_ratio]

    #get the rows of all the close matches in our dataframe
    rows_with_matches = df[column].isin(close_matches)

    # replace all rows with close matches with the input matches
    df.loc[rows_with_matches,column] = string_to_match
    print("Done all")

```

```

[70]: replace_matches_in_column(df =
↳professors,column='Country',string_to_match='south korea')

```

Done all

- 1) And now let's check the unique values in our "Country" column again and make sure we've tidied up "south korea" correctly.

```

[71]: countries = professors['Country'].unique()

```

```

[72]: countries

```

```

[72]: array(['thailand', 'pakistan', 'germany', 'austria', 'australia', 'uk',
        'china', 'france', 'usofa', 'south korea', 'malaysia', 'sweden',
        'italy', 'canada', 'norway', 'ireland', 'new zealand', 'urbana',
        'portugal', 'russian federation', 'usa', 'finland', 'netherland',
        'greece', 'turkey', 'macau', 'singapore', 'spain', 'japan',
        'hongkong', 'saudi arabia', 'mauritius', 'scotland'], dtype=object)

```

```

[ ]:

```

1 THANKYOU KAGGLE

```

[ ]:

```